

3244-2010-0014 - The Pawbability of Adogtion

Authors: Chua Xin Xuan, Lim Xi Chen Terry, Lim York Tee Gorden, Luk Chi Heng, Tan Jie Yi, Tan Yoong Kang Colin

Email: e0425690@u.nus.edu, e0406494@u.nus.edu, e0309693@u.nus.edu, e0310461@u.nus.edu, e0426306@u.nus.edu, e0325403@u.nus.edu

Abstract

In this report, we will aim to predict the adoptability of a pet in an animal shelter given its portfolio on a web profile page. Additionally, we want to identify which traits & profile pictures are desirable when pet adopters choose to adopt a pet. This will be done via the use of Decision Trees and a Convolutional Neural Network models to deal with the complex mix of variable types in our data. Based on our prediction models, we will examine the significant variables and derive insights on the possible actions following our results.

1 Introduction

Some animals seem to be able to get adopted easily into their forever homes, while others only have bleak futures ahead of them as they continue to remain in the shelter.

With the growing popularity of online adoption drives¹, we wish to build a prediction model to predict if an animal would get adopted based on its online portfolio.

Furthermore, to understand this phenomenon better, we also identify which attributes in a pet's web profile, such as gender, age, play a more significant role in determining whether it gets adopted or not. We will also be looking into the profile pictures of the online portfolios to see if there is any correlation between profile pictures of animals that do or do not get adopted.

For this report, we will make use of Decision Tree and Convolutional Neural Network models to try to provide an accurate prediction as well as interpret the results to discern and understand the significant attributes and their behaviour.

Through this, we hope to improve the chances of these animals being adopted, which will help increase the shelters' ability to take in more stray animals, decrease the number of animals being euthanized in shelters or come up with preemptive solutions to cater to animals identified to have trouble in getting adopted.

This report will explain the machine learning methods and decisions we have employed, our subsequent predictions and how we interpreted the findings. For each of these methods, we will first elaborate on its methodology, then evaluate how accurate the predictions generated are based on both real life observations as well as through preliminary exploratory data analysis.

2 Related Work

The dataset we will be working with comes from Petfinder, a non-profit organization. The dataset used in particular comes from the Kaggle Competition² hosted by Petfinder Malaysia in 2018 with the intention of finding the speed in which an animal gets adopted. For the purposes of this report, we only consider animals who do or do not get adopted.

There already exists several studies and reports on the dataset that also choose to focus on both building models to determine a pet's adoption as well as identifying the significant attributes. Such reports include Zhang's Predicting Adoption Speed for Petfinder³, and the various public submissions from the Kaggle Competition.

However, most of these analyses choose to utilise a full neural network deep learning model. Such examples include the best performing Kaggle submission by the team Wodori⁴, prioritising accuracy but leaving little to none interpretability of the results due to the 'black-box' nature of neural network models.

Such works also fail to conduct further testing to determine the significance or understand the correlations between adopted and unadopted pet profiles.

Others, like Zhang, instead choose to only focus solely on the tabular data by modeling them after classical machine learning algorithms, which provides useful insights and interpretation of the data. However, they leave out attributes which we consider are potential key factors that impact the decision behind the adoption of a pet (such as the pictures uploaded).

Taking the above shortcomings into mind, we have decided to find the halfway point between the two extreme directions others have taken. We hope that by employing both highly interpretable machine learning algorithms as well as deep learning neural networks with subsequent testing. We will be able to provide not only accurate results, but a more holistic and complete interpretation of the findings using the data we have available to us, as compared to prior works.

3. Main Methodology

Our dataset comprises a plethora of variable types, notably photos and tabular data regarding characteristics of the pet for adoption.

As such, our initial plan was to build a multi-input model, where we feed the tabular data into a Dense network and image data into a CNN network then merge them and pass it into a common dense layer.




Fig 1: Multi-input model displaying insertion of both tabular data and images into their respective connections into a merging model to generate a prediction

However, due to the 'black-box' nature of deep learning neural

¹Lock, C, (2020), Covid-19 heroes: Animals find fur-ever homes with online adoption drives, <https://www.straitstimes.com/lifestyle/animals-find-fur-ever-homes-with-online-adoption-drives>

² <https://www.kaggle.com/c/petfinder-adoption-prediction/overview>

³ Zhang, E.Z, (2020), Predicting Adoption Speed for Petfinder, <https://towardsdatascience.com/predicting-adoption-speed-for-petfinder-bb4d5befb78c>

⁴ <https://www.kaggle.com/naka2ka/stack-480-speedup-groupfold-with-no-dict>

networks, we will not be able to properly or accurately discern what are the 'significant' variables and their level of importance using such an approach.

Hence, we decided to have two separate approaches:

Approach 1: Decision Trees (DT)

Decision Trees will be used to track and model the mix of categorical, continuous and ordinal features. Given the structure of our data cleaning, transformation and reduction and mix of variable types, we have tried implementing different classification methods as well, which are also included in our code.

These models include Multinomial Naive Bayes, Logistic Regression and K-Nearest Neighbors. We would use them as our baseline models (by comparing their workability, accuracy and interpretation) and compare them in our subsequent final model for tabular data during our evaluation.

In the end, we ultimately chose to use Decision Trees as not only does it give a high enough accuracy score, but we are able to easily interpret and determine important variables by simply observing the nodes and branches of the tree. To corroborate our decision, we performed Decision Trees on our overall dataset to verify if it is implementable, feasible and effective, before doing our split analysis on cats and dogs.

We also considered using a Random Forest as well which would produce a better accuracy score at the cost of interpretability. However, as we are much more interested in the significance of our variables, we valued the high interpretability of the decision tree over a better prediction score.

Approach 2: Computer Vision (CV)

We will use Computer Vision to handle the important spatial informational aspects in the images of animals in our data set. We believe that a Convolution Neural Network is a good choice due to its high accuracy in image classification. Images have high dimensionality, thus CNN will serve to reduce the number of parameters required, while preserving important features. This can decrease the problem of overfitting which will lead to high test cost. Hence, resulting in a high accuracy for image classification.

We had also converted the AdoptionSpeed column to an Adoption column that contains only binary values. Adoption speed of 0 - 3 is classified as Adopted, while adoption speed of 4 is classified as not Adopted.

| Original (Adoption Speed) | Modified (Adoption) |
|--|---------------------|
| 0 - Pet was adopted on the same day as it was listed. | Yes |
| 1 - Pet was adopted between 1 and 7 days (1st week) after being listed. | |
| 2 - Pet was adopted between 8 and 30 days (1st month) after being listed. | |
| 3 - Pet was adopted between 31 and 90 days (2nd & 3rd month) after being listed. | |
| 4 - No adoption after 100 days of being listed. (There are no pets in this dataset that waited between 90 and 100 days). | No |

Fig 2: Table documenting our relabelling from the original dataset (Adoption

Speed) to simply Adoption

4 Decision Trees

4.1 Data Manipulation

To deal with categorical, ordinal and continuous data, we want to manipulate our inputs into a categorical or ordinal format to aid the decision making in our decision tree. Hence, we converted some continuous variables to ordinal data based on their respective quantiles. Columns on "Breed1" and "Breed2" have high dimensions as they can take around 300 different values. Rather than having over 300 more variables from one-hot encoding, we reduced the dimensions by collapsing this into a single binary variable: 1 indicates the animal is a pure breed while 0 indicates its a mixed breed. Additionally, we dealt with "Age" specially, as the difference in months at earlier stages bear more significance than during old age. Hence, a more intricate interval is used to bin "Age" into ordinal data, where we split between [0-1], (1-2], (2-3], (3-6], (6-12], and (12 and above), to account for younger animals in greater detail.

The data set provided only consists of cats and dogs, and the distribution is shown as below:




Fig 3: Breakdown of class distribution in our revised dataset

4.2 Decision Trees Modelling

After cleaning the data, we decided to construct two separate Decision Trees, with prior information that the important features for cats and dogs may differ. We then created a general flow in constructing each of these trees. Firstly, we tried two criterions: Gini Index and Entropy. Secondly, we iterated and plotted accuracy scores across an interval of lengths for maximum tree depth (max_depth). Thirdly, we Identified best max_depth and reviewed accuracy scores near that depth for pruning. In the end, we finalize the choice of max_depth and criterion, then implement Decision Tree for the dataset. With the tree constructed, we can output Accuracy Scores and compare across [Accuracy, Specificity, Sensitivity, F1-Score and AUC], where we tuned with a validation set and used a test set to showcase final scores.

Finally, once we obtain a decision of suitable accuracy, we can identify the important attributes in the decision making process. This is done via "model.feature_importances_", which is based on Gini importance, where how much "impurity" this feature reduces is computed and normalized. The greater the reduction, the better.

4.2.1 Decision Trees for Cats

Having prior belief that cats and dogs may have different important attributes, we split our analysis here and created Decision Trees for each type. Following our modelling above, we chose the criterion to be 'Entropy' and set the 'Max depth' at 6. This is because the optimal accuracy score was achieved at max depth = 6. We went with it as its score is significantly larger than other max_depth

values and has enough representation power while being able to specialize and determine more important features.

4.2.2 Evaluation for Cats (DT)

| Evaluation Metric | Score (To 3 significant figures) |
|-------------------|----------------------------------|
| Accuracy Score | 0.632 |
| Specificity | 0.566 |
| Sensitivity | 0.692 |
| F1-Score | 0.664 |
| AUC | 0.629 |

Fig 4: Evaluation Metrics with the Cat Dataset

4.2.3 Findings for Cats (DT)

For cats, important features were Age, Amount of Photos, Fur Length, Sterilized, Maturity Size, in descending order. Age is the most pivotal feature with the largest Gini importance of 0.579 compared to PhotoAmt at 0.152. Notably, the relationship between MaturitySize and the adoptability of cats is in contrast with dogs, which will be specified later.

Maturity Size for Cats:

(1 = Small, 2 = Medium, 3 = Large, 4 = Extra Large)




Fig 5: Barchart of the proportion of adopted and unadopted cats against their Maturity Size. Blue bar depicts "not adopted", orange depicts "adopted"

Cats of the smallest size are heavily favoured, but sizes larger than small have similar chances of being adopted.

4.2.4 Decision Trees for Dogs

Similarly for dogs, we chose the criterion to be 'Entropy' and set the 'Max_depth' to be 8. This was because the optimal accuracy score was achieved at 'Max_depth' = 8. We went with it as once again, its score is significantly larger than other max_depth values and has enough representation power while being able to specialize and determine more important features.

4.2.5 Evaluation for Dogs (DT)

| Evaluation Metric | Score (To 3 significant figures) |
|-------------------|----------------------------------|
|-------------------|----------------------------------|

| | |
|----------------|-------|
| Accuracy Score | 0.608 |
| Specificity | 0.659 |
| Sensitivity | 0.552 |
| F1-Score | 0.574 |
| AUC | 0.606 |

Fig 6: Evaluation Metrics with the Dog Dataset

4.2.6 Findings for Dogs (DT)

For dogs, the important features were Age, Amount of Photos, Maturity Size, Fur Length, Sterilized. Despite having the same attributes as cats, the importance across these features are rather equal, having similar values for Gini importance where Age is only 0.07 higher than Amount of Photos. Previously mentioned, the relationship for Maturity Size and adoptability of dogs differs in distribution from cats, hence we will showcase the contrast.

Maturity Size for Dogs:

(1 = Small, 2 = Medium, 3 = Large, 4 = Extra Large)




Fig 7: Bar chart of the proportion of adopted and unadopted dogs against their maturity size.

Dogs of the largest size are heavily favoured, but sizes larger than that have similar adoption chances of being adopted.

4.3 Overall Discussion

With both our Decision Trees pruned and finalized, we can now evaluate our model across the different baseline models mentioned in Section 3. We can compare the accuracy scores below:

| Classifier | Score | Decision Tree | Score |
|-------------------------|-------|---------------|-------|
| Multinomial Naive Bayes | 0.602 | Cats | 0.632 |
| K-Nearest Neighbors | 0.518 | Dogs | 0.608 |
| Logistic Regression | 0.628 | | |
| Decision Forest | 0.639 | | |

Fig 8: Accuracy Scores using Baseline Models

Our two Decision Trees performed generally well in terms of prediction, despite being slightly lower in scores than some baseline models. However, we note that these baseline models are difficult to interpret which variables bear more weight, despite being fast in implementation. Therefore, since our decision trees have similar performance in terms of accuracy and are able to determine feature importance easily, it was optimal to use Decision Trees as our final model; even though it may require more meticulous details in structuring, where pruning and some dimensionality reduction was conducted.

After evaluating, we would like to consolidate our findings from both Decision Trees. Besides Maturity Size, all other important attributes were common in both Cats and Dogs. We note that they vary in the order of importance, where some attributes bear more weight depending on the type of animal. We analyzed their distributions separately for each of these features and realized that they are very identical. Hence, a combined evaluation is depicted below, to reduce repetition.

Age:

As the pet grows older, its chances of adoption severely decreases, especially in the earlier months.




Fig 9: Bar chart of the proportion of adopted and unadopted dogs against their age.

Fur Length:

Longer Fur Length seems to be more appealing, as seen by the percentage of adoption increasing as FurLength increases. This could imply that fluffy and furry pets are more attractive as compared to bare-skinned ones.




Fig 10: Bar chart of the proportion of adopted and unadopted dogs against their Fur Length

Sterilized:

(Binary, 1 = Has been sterilized, 0 = Has not been sterilized):

Surprisingly, pets who are sterilized have a much lower chance of being adopted.




Fig 11: Bar chart of the proportion of adopted and unadopted dogs against if they have been sterilized

Amount of Photos:

Surprisingly, pets with a smaller amount of photos tend to be adopted more often.




Fig 12: Bar chart of the proportion of adopted and unadopted dogs against their number of photos.

However, we note that when pets have no photos at all, their chances of adoption are much lower, emphasizing the importance of our Computer Vision model to derive more significant insights.

Concluding our segment on dealing with tabular data through Decision Trees, our next section would elaborate on our Computer Vision model in detail, in order to determine the significant factors in the images of the pets.

5 Computer Vision

5.1 Method

Data exploration and analysis was first carried out to prepare the data for training.

1. Rows containing no images were dropped, because this model requires images as data instead of tabular data.

2. The first image for pets that have more than one photo was used for the model, as to ensure fairness when comparing between each pet. Furthermore, since the first image is the profile picture of the pet in the Adoption webpage, it will most likely have the greatest impact on the potential pet adopters when they first see an animal.

3. For the percentage of data to be used for training, testing and validation, 80% of the data was first used for training and 20% for testing. However, we realised that we were unable to use the Kaggle platform for testing as the problem was simplified from a multi-class

problem to a binary-class problem instead. Thus, the bottom 10% of the data was used for testing instead. Out of the remaining 90%, 80% was used as training and 20% for validation.




Fig 13: Train Test Split of Dataset

4. Similar to decision trees, we have also split the data into cats and dogs, because we believe that the features of each animal are different when attempting to predict their adoptability. Potential pet adopters may also be looking out for different features for different pets. For example, dog adopters may be looking out for dogs which are bigger in size and darker coated, while for cats, they may be looking out for cats which are smaller in size and lighter-coated.

5.2 Model

Model: Pre-trained Xception model + Additional hidden layers
 Activation function: Sigmoid
 Optimizer: Stochastic Gradient Descent
 Loss function: Binary Crossentropy
 Metric: Accuracy

| Model: "sequential" | | |
|----------------------------------|--------------------|----------|
| Layer (type) | Output Shape | Param # |
| xception (Functional) | (None, 8, 8, 2048) | 20861480 |
| global_average_pooling2d (Gl | (None, 2048) | 0 |
| output (Dense) | (None, 1) | 2049 |
| Total params: 20,863,529 | | |
| Trainable params: 2,049 | | |
| Non-trainable params: 20,861,480 | | |

5.3 Evaluation

| Model | Accuracy of Dog Model (To 3 significant figures) | Accuracy of Cat Model (To 3 significant figures) |
|--|---|---|
| Simple CNN Model | 70.2% | 72.5% |
| DenseNet121 | 65.0% | 73.1% |
| InceptionV3 Model | 70.5% | 73.1% |
| Xception Model (w/o Data Augmentation) | 72.0% | 73.0% |
| Xception Model (w Data Augmentation) | 72.5% | 73.5% |

Fig 14: Accuracy Scores of Models

Our baseline model was a simple image classification model built from scratch. However, we realised that if we wanted to further improve its accuracy, it would take us a longer period of time and more resources to fine tune the model. Hence, we had decided to

utilise transfer learning⁵ which is to use a pre-trained model⁶ with its trained weights as the base model. A pre-trained model has a good representation of features for over a million images belonging to a thousand different categories. It can act as a good feature extractor for new images making it suitable for computer vision problems.

Among all the various pre-trained models that we have tested, we ultimately chose to use Xception as it has one of the highest accuracy. Furthermore, since it has been previously trained on images for image classification tasks, this makes it a good fit for our task.

As our model requires a probability value as an output to decide whether a pet is adopted or not, we have narrowed down to either Softmax or Sigmoid among all the various activation functions. Comparing between these two, Softmax is used for multiclass problems, while Sigmoid is used for 2-class problems. Hence, sigmoid is the most suitable activation function to be used in our model.




Fig 15: Graph of Training and Validation Accuracy & Loss for Xception Model (w/o Data Augmentation)




Fig 16: Graph of Training and Validation Accuracy & Loss for Xception Model (w Data Augmentation)

After running our model and plotting the trained loss and validation loss for each epoch, we observed that there was a large difference

⁵ Dipanjan(DJ) Sarkar, (2018), *A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning*, <https://towardsdatascience.com/a-comprehensive-hands-on-guide-to-transfer-learning-with-real-world-applications-in-deep-learning-212bf3b2f27a>

⁶ Billy Lamberta, Bastian Eichenberger and others, (2017), *Transfer learning and fine-tuning*, https://www.tensorflow.org/tutorials/images/transfer_learning

between the 2 losses due to overfitting as shown in Fig 15. This might be because the amount of data which we can work with has shrunk significantly after filtering and splitting the data into dogs and cats. Thus with a smaller dataset for the model, the model kept seeing the same instance over time across each epoch, resulting in overfitting. Hence, we have used image augmentation whereby we applied some image transformation operations to the dataset (e.g. rotation, shearing, translation, zooming, etc.) to produce new, altered versions of existing images. Therefore, this would artificially increase the size of our training data while maintaining the integrity of the data.

5.4 Discussion

Possible reasons for pets not being adopted include:

1. Dilution of images with text and other images
2. Collaged images
3. Background noise (e.g. presence of human)
4. Irrelevant images
5. More than one pet in the photo

Some examples are as shown below:



Fig 17A: Image with a lot of dilution of other irrelevant images and text



Fig 17B: Image with where a girl is present (Background Noise)

We feel that the profile photo which is the first image of every pet is crucial as it is the photo that gives potential adopters their first impression. This affects whether a potential adopter would click in to find out more about the pet.

To help boost a pet's chances of getting adopted, we believe that the pet should be the main focus in its profile photo. It would also yield higher adoptability for the pet if it were to look directly at the camera.

6 Conclusion

6.1 Summary

Although both models have an accuracy score of approximately 60%, which performs a bit below our expectations. Our research is still useful as our findings suggest that pet adopters place the highest importance on the quality of photos and age of the animal. Higher quality photos and younger pets tend to yield a much higher chance of being adopted. We also found that furry animals are generally more well liked and maturity size plays a vital role. Therefore, in order to improve the adoptability of a pet, it is more effective to work on these key details, disregarding other noisy attributes that may seem important but do not actually play a large role. This would save time, costs and resources, ultimately improving the overall adoption rate.

6.2 Shortcomings

We understand that many other important external factors such as location of shelter, number of volunteers and geographical location were not taken into account in our project due to the lack of data.

Our decision tree could have used more extensive pruning on with regards to other hyperparameters⁷ such as min_impurity_decrease, min_impurity_split, min_samples_leaf, max_leaf_nodes and max_features as opposed to just pruning based on max depth.

Furthermore, given enough time, we could also implement bounding boxes in the images to pinpoint the location of the animal to allow for better identification of the animal.

A possible ethical issue that may arise is the possibility that sheltered animals may receive an unequal level of treatment based on the model's findings. Shelters may choose to allocate more resources on animal's predicted to not be adopted. In the event that the prediction is misclassified, animals already struggling to get adopted may face a tougher time.

6.3 Future Work

Overall, with all the important features identified, we hope that pet shelters make use of our findings to make guided and focused improvements of the profile of the animals in their shelters, thereby increasing the chances of these animals finding their forever homes.

We wish to follow up and disseminate our research and findings to animal protection organisations in Singapore such as SPCA, SOSD and ACRES. We hope that we can apply such methods of Machine Learning and research into Singapore's context to better understand the local landscape of the pet adoption culture.

⁷ Sodner Yildirim, (2020), *Hyperparameters of Decision Trees Explained with Visualizations*, <https://towardsdatascience.com/hyperparameters-of-decision-trees-explained-with-visualizations-1a6ef2f67edf>

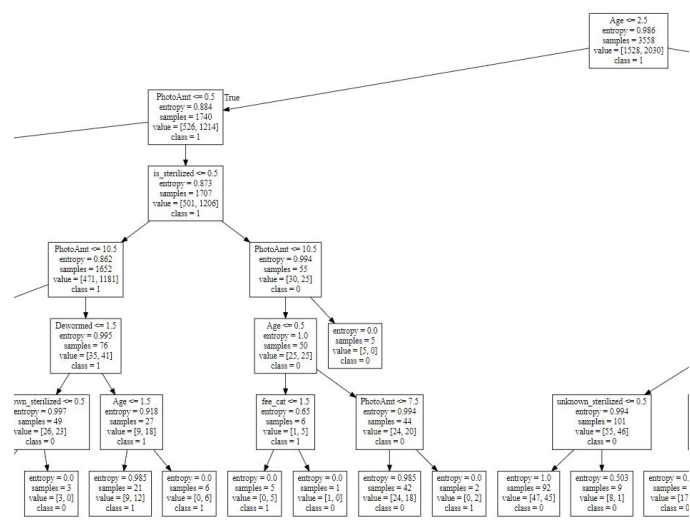
References

- [1] Lock, C, (2020), *Covid-19 heroes: Animals find fur-ever homes with online adoption drives*,
<https://www.straitstimes.com/lifestyle/animals-find-fur-ever-homes-with-online-adoption-drives>
- [2] Kaggle, *PetFinder.my Adoption Prediction*,
<https://www.kaggle.com/c/petfinder-adoption-prediction/overview>
- [3] Zhang, E.Z, (2020), *Predicting Adoption Speed for Petfinder*,
<https://towardsdatascience.com/predicting-adoption-speed-for-pet-finder-bb4d5befb78c>
- [4][kaggler-ja] Wodori,(2019), *stack 480 + speedup + GroupKFold with no dict*,
<https://www.kaggle.com/naka2ka/stack-480-speedup-groupkfold-with-no-dict>
- [5] Dipanjan(DJ) Sarkar, (2018), *A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning*,
<https://towardsdatascience.com/a-comprehensive-hands-on-guide-to-transfer-learning-with-real-world-applications-in-deep-learning-212bf3b2f27a>
- [6]Billy Lamberta, Bastian Eichenberger and others, (2017), *Transfer learning and fine-tuning*,
https://www.tensorflow.org/tutorials/images/transfer_learning
- [7]Sodner Yildirim, (2020), *Hyperparameters of Decision Trees Explained with Visualizations*,
<https://towardsdatascience.com/hyperparameters-of-decision-trees-explained-with-visualizations-1a6ef2f67edf>

Appendix A

<https://github.com/CS3244-2010-0014/Pawbability-of-adogtion>

Portion of Decision Tree for Cats:



As the tree is too big to be shown, please refer to project github repository, cat.dot and dog.dot

